



# Microsoft's Approach to Agentic AI

Framework, Tools, and Best Practices

**Agentic AI** refers to AI systems (or agents) that can **perceive, reason, and act autonomously** toward goals with minimal human intervention<sup>1 2</sup>. Microsoft's overall approach to agentic AI centers on **empowering organizations to start small, then scale up** to more advanced AI agents, all while maintaining **strong governance, transparency, and alignment with business objectives**. In practice, this means **using the right tools for the right stage** of maturity - from quick no-code copilots to fully custom AI agent implementations - and embedding responsible AI principles at each step<sup>3 4</sup>. The core idea is to unlock automation beyond chatbots: not just answering questions, but **completing tasks and orchestrating workflows** for real business impact<sup>5</sup>. Microsoft provides a **framework to guide customers** in choosing appropriate agentic AI tools and techniques based on their needs, and offers an integrated tech stack (like **Azure AI Foundry** and **Copilot Studio**) to support these solutions. Crucially, Microsoft advises enterprises to adopt agentic AI **progressively and responsibly** - ensuring **clear goals, robust safeguards, and measurable outcomes** from the start.

## Phased Framework for Agentic AI Adoption

Microsoft recommends a **phased adoption framework** for agentic AI, which helps organizations match the **complexity of tools and techniques to their business needs**. This framework has three levels of increasing capability and customization, often described as **No-Code**, **Low-Code**, and **Pro-Code** stages<sup>6 7</sup>. At each stage, different Microsoft solutions are available to build AI agents, targeting different users and use cases. The table below summarizes the **tiers and their characteristics**:

Approach & Intended Platform	Intended Users	Suitable Scenarios (Examples)	Capabilities & Integration	Time to Value

<b>No-Code: Microsoft 365 Copilot - Agent Builder (built into M365)</b>	Business users ("smart users", team leads) <sup>8</sup>	Quick wins: FAQ bots, personal assistants, simple task helpers (e.g. answer policy questions or draft emails) <sup>9</sup>	<b>Integrated</b> in everyday tools (Teams/Outlook); connect to internal content (SharePoint, Docs) and perform basic actions (create task, send summary) <sup>10</sup> . Minimal setup, uses org's existing M365 data.	<b>Fastest</b> - build within days; point-and-click configuration <sup>11</sup> .
<b>Low-Code: Copilot Studio (Power Virtual Agents evolution)</b>	"Citizen developers", power users <sup>12</sup>	Intermediate: Multi-turn chatbots, departmental virtual agents, workflow automations with conditional logic (e.g. an HR bot that answers queries and files tickets) <sup>13</sup> <sup>14</sup>	<b>Extensible</b> via connectors and APIs; design dialog flows, call internal/external APIs, apply business rules, and orchestrate multi-step conversations without full coding <sup>15</sup> . More custom logic and enterprise integration (CRM, HR systems, etc.) than no-code.	<b>Moderate</b> - deploy in weeks; requires planning conversations and connectors <sup>16</sup> <sup>17</sup> .
<b>Pro-Code: Azure AI Foundry - Agent Service &amp; SDK</b>	Developers, IT architects <sup>18</sup>	Advanced: Mission-critical AI agents and multi-agent systems needing custom models, complex reasoning, or high reliability (e.g. a	<b>Fully customizable</b> and <b>enterprise-grade</b> : choose or fine-tune LLM models (Azure OpenAI GPT-4, Llama, etc.) <sup>21</sup> ; compose multiple agents with specialized roles; server-side tool orchestration, robust debugging, CI/CD integration, monitoring dashboards <sup>22</sup> <sup>23</sup> . Integrates with identity (Entra ID), enterprise	<b>Longer</b> - project lifecycle (several weeks to months) for development and testing <sup>24</sup> , but yields production-ready,

		<p>financial research agent that analyzes reports &amp; executes trades with approval) <sup>19 20</sup></p>	<p>data stores, and ensures security, compliance &amp; scalability.</p>	<p>governed solutions.</p>
--	--	---	---	----------------------------

**No-Code (M365 Copilot - Agent Builder):** At the entry level, Microsoft enables business users to create simple AI agents *within the Microsoft 365 environment* using natural language, without writing any code. With the **M365 Copilot's Agent Builder**, a user can quickly turn a prompt or a document into a chat-based assistant inside tools like Teams or Outlook<sup>25 26</sup>. This approach is ideal for **fast prototyping and quick wins** – for example, an employee could spin up an FAQ bot for common policy questions or a personal “copilot” that drafts routine communications. Because it’s integrated into M365, the agent can be easily grounded on internal SharePoint knowledge or emails and even perform simple actions like sending a meeting reminder or creating a to-do in Planner<sup>27</sup>. Microsoft advises keeping these no-code agents **very focused in scope** (e.g. “one job to be done”) and leveraging a single high-quality knowledge source for grounding<sup>28</sup>. By piloting with a small group and measuring outcomes (like what percentage of queries it resolves, or how much time it saves), teams can validate value quickly<sup>29 30</sup>. *Guardrails:* Even at this stage, Microsoft emphasizes setting clear expectations for the agent – document its **mandate and limitations** (what it will and won’t do) so users trust it and know its boundaries<sup>31</sup>. Keeping the content curated and responses concise helps maintain quality and user trust from day one<sup>32</sup>.

**Low-Code (Copilot Studio):** As needs grow, Microsoft encourages moving to low-code solutions via **Copilot Studio**, which evolves from the Power Platform (and Power Virtual Agents) to allow more complex conversations and integrations<sup>33 34</sup>. In Copilot Studio, power users or IT pros can design multi-turn dialog flows with branching logic, incorporate **custom connectors** to third-party APIs or databases, and embed the agent into business workflows without building everything from scratch<sup>35</sup>. This is useful for departmental or cross-functional bots – for example, a **customer service agent** that not only answers questions but also creates a support ticket in ServiceNow, or an **HR onboarding agent** that guides a new hire through tasks and triggers IT

provisioning workflows. Copilot Studio provides a visual conversation designer and supports calling external systems securely (with connectors and permissions), along with testing tools and analytics<sup>36 37</sup>. It effectively adds an orchestration layer: you can structure the conversation (e.g. greet → clarify user intent → retrieve info or perform action → confirm → summarize) and implement **business rules** (like validating inputs or requiring manager approval for certain actions)<sup>38</sup>. Microsoft notes this approach can handle more **extensive automation** while still being accessible to a broader range of builders. Typical **time-to-value** is on the order of a few weeks – setting up connectors, refining dialogues, and user acceptance testing<sup>39</sup>. Best practices: Microsoft recommends **separating knowledge from logic** – keep your content sources where they’re governed (e.g. SharePoint for documents) and use Copilot Studio to manage the conversation and process flow<sup>40</sup>. It’s also important to **instrument outcomes** (not just chat metrics): track whether the agent successfully completed tasks (like how many tickets opened) to gauge its real impact<sup>41</sup>. Low-code agents should enforce basic validation and have fallback paths (e.g. escalate to a human or a supervisor) when confidence is low<sup>42</sup>.

**Pro-Code (Azure AI Foundry - Agent Service):** For **enterprise-grade agentic AI applications**, Microsoft provides Azure AI Foundry – essentially an “**AI agent factory**” that gives developers full control to build, test, and deploy autonomous agents at scale<sup>43 44</sup>. Azure AI Foundry’s Agent Service is a managed runtime that ties together **models, tools, and governance** in one system<sup>45</sup>. With Foundry, developers can pick the optimal **LLM model** (from Azure OpenAI GPT-4 to open-source models like Llama) for each agent’s reasoning<sup>46</sup>, and customize it via fine-tuning or prompt engineering to imbue domain knowledge and desired behavior<sup>47 48</sup>. Crucially, agents in Foundry are **tool-using and goal-driven** – you explicitly equip them with APIs or plugins so they can take actions (e.g. call a database, invoke an Azure Function) rather than just chat<sup>49</sup>. Foundry supports building **multi-agent systems** too: agents can collaborate or message each other to solve parts of a problem, coordinated by an orchestration layer<sup>50 51</sup>. All of this runs within a **secure, observable, and governable runtime** that Azure manages<sup>52</sup>. For example, the platform logs every agent decision and tool invocation as a structured “thread” (conversation log), giving developers and auditors full visibility into what the agent is doing<sup>53 54</sup>. Policies for **content filtering and safety** are built-in – preventing certain unsafe outputs or flagging them – so the agent operates within corporate compliance boundaries<sup>55 56</sup>. Foundry also integrates with **Microsoft Entra ID (Azure AD)** for identity and Role-Based Access Control, meaning each agent can have a scoped identity with permissions, just like a human employee or service account<sup>57 58</sup>. This is key when agents are able to execute transactions or modify records: you manage what they’re allowed to do. The trade-off

with the pro-code approach is a **longer development cycle** and a need for software engineering skills - you will write code (using the Foundry SDK/CLI) to define agent logic, and you'll apply rigorous testing (including offline evaluation with test datasets, and continuous monitoring in production)<sup>59 60</sup>. The benefit is **maximum flexibility and reliability**: you can choose specialized model strategies, implement complex reasoning patterns (e.g. **Tool-Reflex-Plan loops, or ReAct frameworks**), and optimize for performance and cost at a granular level<sup>61 62</sup>. Microsoft positions Azure AI Foundry as the path to go from promising prototypes to **production-ready AI workflows** in the enterprise, by providing the necessary scaffolding (or "**unified agent platform**") so teams **don't have to reinvent infrastructure** for orchestration, logging, security, etc.<sup>63 64</sup>.

**Choosing the Right Approach:** The intent of this tiered framework is to guide organizations **where to start and when to level up**. Microsoft often advises: **start with the simplest solution that delivers value**, and only graduate to more complex builds as needed<sup>65 66</sup>. Many business scenarios can begin with a **no-code pilot** to prove the concept (for example, an internal knowledge base bot that answers employee questions) - this can be done in days and demonstrates the potential time saved or improved response rate<sup>67</sup>. If that pilot's scope needs to expand (say, the bot needs to perform **transactions** or handle multi-turn troubleshooting dialogs), it might be time to move into Copilot Studio for a more feature-rich agent with moderate development effort<sup>68 69</sup>. Finally, for use cases that demand **high correctness, custom integration, or scale** - e.g. an agent that handles financial data or mission-critical processes - the pro-code route with Azure AI Foundry is appropriate<sup>70 71</sup>. Microsoft's **guidance to customers** is to avoid over-engineering too early: "**stay on course and go progressive**", in the words of one Microsoft architect<sup>72</sup>. This phased approach accelerates adoption (by capturing quick wins) while laying a path to robustness; it also helps organizations **build internal skills step by step** - from tinkering with no-code agents to eventually managing sophisticated AI systems.

Above is an example **90-day roadmap** Microsoft suggests for rolling out agentic AI in a controlled manner<sup>73 74</sup>. By starting with rapid prototypes and progressively adding complexity (and oversight), organizations can **learn and adapt** before fully operationalizing an AI agent<sup>75</sup>. This staged approach echoes best practices from software development, ensuring that by the time an agent is serving critical business needs, it has been vetted for quality, safety, and ROI.

Microsoft's Agentic AI Tech Stack in Action

Microsoft's technology stack for agentic AI provides the **building blocks for each stage** of the above framework – from integrated 365 Copilots up to Azure cloud services. Below we highlight two key offerings, **Copilot Studio** and **Azure AI Foundry**, and how they support agentic AI implementations with real-world examples:

- **Microsoft 365 Copilot & Copilot Studio:** For organizations using Microsoft 365, **Copilot Studio** (part of the Copilot suite) is the primary tool to create and manage AI agents without coding. It leverages the groundwork of the **Power Platform**, meaning it inherits robust connectors and a friendly interface for building conversational bots and automation flows<sup>76</sup>. In Copilot Studio, a user can define an agent's knowledge sources (e.g. specific SharePoint sites), conversation topics, and integrate actions like updating a record in Dynamics 365 or sending an email via Outlook. All of this is done through configuration and natural language prompts – effectively a *low-code chatbot builder*. A recent update even introduced an "**Agent Store**" and Copilot **Agent Builder lite** accessible directly to business users, showing Microsoft's push to democratize agent creation. For example, **Microsoft 365 Copilot Workflows** allow users to automate multi-step tasks by simply describing them (e.g. a bug reporting workflow that logs an issue and notifies the team). Under the hood, Copilot Studio handles converting these descriptions into working Power Automate flows or bot logic. This tight integration means companies can quickly tailor Copilot to their unique processes. Microsoft's recommendation is to use Copilot Studio for **departmental agents** that require some custom logic but *not* extensive coding – cases like a **sales assistant bot** that can look up client info and prepare a briefing, or a **workflow agent** that coordinates meeting scheduling across calendars. One reason Copilot Studio has seen rapid uptake (over 230,000 organizations are already using it, including 90% of Fortune 500<sup>77</sup>) is that it provides immediate value on top of existing M365 investments and governance. Data created or used by these agents stays within the Microsoft 365 environment, respecting the same compliance and DLP (Data Loss Prevention) policies that already protect emails, documents, etc.<sup>78</sup>. In short, Copilot Studio allows companies to **scale out AI-assisted automation safely within their tenant**, before they venture into custom AI development.
- **Azure AI Foundry (Agent Service):** Azure AI Foundry is Microsoft's **enterprise agent development platform**, combining the power of Azure's AI models with a managed runtime purpose-built for autonomous agents<sup>79</sup>. In practice,

Foundry is what customers use when they need an agent (or **multiple cooperating agents**) to perform sophisticated tasks, reliably and under governance. For example, consider a scenario in financial services: a bank wants an AI agent to **analyze market reports** and autonomously execute trades or recommendations according to a strategy. This agent must be highly trustworthy (tested against regulations, with audit logs) and might need to break down goals - read reports, summarize findings, forecast an outcome, then act. Azure AI Foundry is designed for such complex workflows. It provides templates and components so developers can build agents that possess key **"agentic" capabilities: tool use, reflection, planning, and even multi-agent collaboration**<sup>80 81</sup>. For instance, using Foundry a team can create a set of specialized agents (one focused on data gathering, one on analysis, one on execution), and an orchestrator agent that delegates tasks among them - an architecture similar to how enterprises structure human teams<sup>82</sup>. Foundry's Agent Service ensures all these interactions happen in a controlled way: it manages the sequence of tool calls, handles errors or retries, and logs every step in an **Application Insights** telemetry stream for later review<sup>83 84</sup>. Moreover, it comes with **enterprise hardening out-of-the-box**: content **filters to prevent toxic outputs or data leakage**, network isolation options, encryption of agent state, and full **identity integration** so that any action an agent takes can be attributed and access-controlled<sup>85 86</sup>. This directly addresses the challenge many organizations face when moving from a prototype (often just a clever prompt on a GPT model) to a real application - how to implement all the "scaffolding" like security, monitoring, and compliance. With Foundry, Microsoft essentially provides that scaffolding. A **unified platform is critical**, Microsoft notes, because without it, teams might build ad-hoc solutions for chaining model prompts, calling APIs, logging outputs, etc., which can be fragile and hard to govern<sup>87</sup>. By standardizing on Foundry, enterprises get a consistent framework to develop agents across use cases. For example, Microsoft has highlighted how **Fujitsu** used agentic AI (on Azure) to transform its sales proposals process: they built a collection of agents for data analysis, market research, and document generation that worked together to produce complete proposal documents - reducing proposal prep time by **67%**<sup>88</sup>. Another case is **ContraForce's security automation**: they leveraged planning agents (on Azure AI) to automate 80% of cybersecurity incident handling steps, cutting response times and even achieving a cost of under \$1 per incident for processing full investigations<sup>89</sup>. These examples show Azure AI Foundry's strength in orchestrating complex, multi-step workflows that go well beyond a

chat-style Q&A. The platform allowed those companies to plug in custom logic (their own APIs, their own data) and still rely on Microsoft's AI and governance capabilities. It's also noteworthy that Azure AI Foundry supports **advanced model tuning** techniques - e.g. **fine-tuning LLMs with enterprise data, reinforcement learning from human feedback (RLHF), and Direct Preference Optimization (DPO)** - to align the AI agents with company-specific requirements<sup>90 91</sup>. This means a Foundry-built agent can have not just generic GPT-4 intelligence, but one that speaks in the company's tone, respects industry regulations, and continuously learns from feedback in production. All these features reflect Microsoft's overall approach: enabling customers to harness cutting-edge AI (like GPT-4 or emerging open models) *on their terms* - with control over data, iterative improvement, and confidence in safety.

## Responsible Adoption: Governance, Transparency & Alignment

Throughout Microsoft's guidance on agentic AI, a consistent theme is **responsible AI adoption**. Because autonomous agents can make decisions and initiate actions, Microsoft urges organizations to **evolve their governance practices** and ensure each AI agent remains **transparent in operation and aligned with business goals**<sup>92 93</sup>. Key recommendations include:

- **Establish Governance Mindset for AI Agents:** Treat AI agents as a new type of *digital workforce* within the enterprise. Just as you manage human employees or software services, you should **assign identities to agents, define their roles and permissions, and monitor their activities**<sup>94 95</sup>. Microsoft's Work Trend Index (2025) describes leading firms moving from using AI as assistants to deploying "*digital colleagues*" that handle entire business processes<sup>96</sup>. In this paradigm, not every agent gets full autonomy out of the gate. Organizations should **tier agent autonomy levels** - for example, one tier of agents might only answer informational queries, while another tier (higher trust) can execute transactions under oversight<sup>97</sup>. Microsoft likens this to how you wouldn't give a new human hire admin privileges on day one<sup>98</sup>. By defining such **guardrails** (perhaps in policy: which agents can do what), companies prevent overreach. Additionally, Microsoft suggests establishing **human oversight roles** for agents: *Reviewers* to verify AI outputs for accuracy, *Monitors* to track agent actions in real-time, and *Protectors* to intervene or adjust permissions if an agent behaves outside bounds<sup>99</sup>. This

structured supervision ensures agents remain under appropriate check, especially as they start taking on higher-risk tasks.

- **Apply Existing (Low-Code) Governance to Agents:** One advantage for Microsoft customers is that many governance practices used for the Power Platform and low-code tools can **carry over directly to AI agents**<sup>100</sup>. Microsoft advises leveraging your established **Center of Excellence (CoE)** for Power Platform to also oversee AI agent projects<sup>101</sup>. Security measures like **Data Loss Prevention (DLP) policies**, environment management, and **role-based access control** that were set up for Power Apps/Power Automate should be extended to Copilot Studio and other agent builders<sup>102</sup>. For example, if a company has blocked certain sensitive data from being used in Power Platform apps, the same rules should apply to agents (so an AI agent can't inadvertently access a restricted database). Tools such as **Microsoft Purview** can help track and enforce data governance, and **Entra ID** (Azure AD) ensures that agent identities have appropriate access **scopes**<sup>103</sup>. The message is: *don't reinvent the wheel for agent governance* – reuse the compliance and security frameworks you have, updating them to account for the new behaviors of AI agents. Microsoft has published IT adoption guides and even an "**agents adoption**" site to help administrators integrate these governance steps into their agent rollout<sup>104</sup>.
- **Drive Transparency with Telemetry and Control Costs: Visibility** is the cornerstone of governing autonomous systems<sup>105</sup>. Microsoft emphasizes setting up robust **telemetry and analytics** for every agent in operation<sup>106</sup>. This means tracking who created an agent, what data it's using, how often it's being invoked, and what outcomes it produces<sup>107</sup>. Copilot Studio, for instance, comes with built-in analytics dashboards that show usage stats and performance of agents<sup>108</sup>. Likewise, Azure AI Foundry logs detailed traces (every step of an agent's reasoning and tool usage) which can feed into **Application Insights** or SIEM systems for monitoring<sup>109</sup>. By reviewing these logs, IT teams can **spot anomalies or inefficiencies** – e.g. an agent that is rarely used (maybe it's redundant or needs improvement) or one that's calling an API too often (incurring cost). Microsoft also provides a **cost management toolset** (like an AI cost calculator and Azure cost alerts) to help forecast and cap the expenses associated with running AI agents<sup>110</sup>. The goal is not only to prevent budget surprises but also to ensure each agent is actually delivering business value proportionate to its cost. CIOs are encouraged to ask "what outcomes are agents driving?" rather than just "how much are we spending"<sup>111</sup> – shifting

focus to **business impact**. If an agent isn't pulling its weight (for example, an internal sales assistant that hasn't improved any KPI), the logs and usage data should reveal that, prompting a course correction or retirement of that agent. In summary, **transparency** in agent operations – via logging, audits, and regular performance reviews – is critical to maintain trust and efficacy.

Microsoft's stance is that *governance without visibility is just guesswork*<sup>112</sup>, so they equip their agent platforms with rich monitoring capabilities to avoid the AI "black box" problem.

- **Empower Innovation with Guardrails:** Microsoft encourages organizations to let business users experiment with agentic AI (since those closest to the work often have the best ideas for improvement) – but to do so within a **managed sandbox**<sup>113</sup>. A recommended approach is a "**zoned**" **governance model**<sup>114</sup>:
  - *Zone 1: Personal Productivity* – a safe playground where individual employees can try creating agents for their own tasks in isolated environments<sup>115</sup>. This zone has strict limits (no broad data access, maybe sample data only) to ensure any mistakes are low-impact, and it comes with guiding policies so users learn responsible practices.
  - *Zone 2: Team/Department Collaboration* – a middle tier where agents can be developed for a team's use with stronger controls in place<sup>116</sup>. Here, environment-level policies and connector restrictions are enforced by IT, and there's oversight on what the agent is allowed to do in shared systems. This enables broader adoption (multiple users relying on an agent) while keeping compliance in check.
  - *Zone 3: Enterprise Managed* – the highest tier for **production-grade agents** that serve at scale or perform critical functions<sup>117</sup>. In this zone, agents undergo rigorous review, have **continuous monitoring**, and are subject to formal lifecycle management (versioning, testing, approvals). Security is tight and any autonomous decisions are logged and often require secondary validation. This progressive autonomy model prevents chaos: it ensures that by the time an agent is widely deployed, it has passed through controlled stages of maturity. Microsoft notes that scaling agent deployment is not just about technology but also **organizational structure and roles**<sup>118</sup>. New roles are emerging, such as *AI Product Managers, AI Ethicists, Prompt Engineers*, etc., to steward these projects. By empowering business units to innovate in early zones and then involving specialized teams for later zones, companies can **capture grassroots innovation** while still applying professional rigor when

it counts. The Power Platform's existing environment segmentation (Dev/Test/Prod environments, for example) can be leveraged to implement these zones for agents<sup>119 120</sup>. All agents, no matter the zone, should operate within **strict security and compliance boundaries** – for instance, an agent should only access data sources its role permits, and every external action (like sending an email or modifying a record) should be auditible<sup>121</sup>.

- **Align Agents with Business Goals & Foster Adoption:** Microsoft's guidance stresses that agentic AI should not be pursued as a tech novelty – it must align to real business objectives and have stakeholder buy-in. Practically, this means **defining success metrics upfront** for any agent project (e.g. reduce response time by 50%, or handle 80% of Tier-1 support tickets)<sup>122 123</sup>. By tying the agent's purpose to a KPI, teams ensure the AI remains a tool for business value, not just a fancy demo. Microsoft also recommends a "*measure and iterate*" approach: start with a clearly defined outcome, instrument the agent to capture metrics on that outcome, and iterate based on feedback and results<sup>124</sup>. On the cultural side, driving adoption of agentic AI solutions requires **education and engagement**. Microsoft highlights that the toughest challenges are often cultural, not technical<sup>125</sup>. To overcome this, they suggest building an **internal community of practice** around AI agents<sup>126</sup>. For example, host regular "Agent Show-and-Tell" sessions or hackathons where teams present what their agents do<sup>127</sup>. Recognize and reward successful agent projects and identify **champions** in each department who can mentor others in using or creating AI agents<sup>128</sup>. Additionally, invest in training programs that cover both how to develop agents (for those building them) and how to use them responsibly (for end users making decisions with agent output)<sup>129</sup>. Microsoft provides learning paths (as referenced in their documentation and blogs<sup>130 131</sup>) to help different user personas get up to speed with Copilot, Studio, and Foundry. By fostering a positive, skill-building environment, organizations can **increase adoption steadily** while maintaining oversight. In terms of transparency to users, Microsoft encourages that agents be **identifiable and predictably behaved** – users should know they're interacting with an AI and be informed about the agent's scope. This can be achieved by clearly describing the agent's role when it's introduced and even providing guidance on how to use it (for instance, a chatbot might introduce itself: "I'm a virtual agent that can help with IT support by diagnosing issues or filing tickets. I may redirect complex queries to a human."). Such transparency builds trust and sets the right expectations, aligning well with

Microsoft's Responsible AI principle of **transparency**. It also aligns the agent with business goals by making sure everyone understands the agent's intended value proposition.

**In conclusion**, Microsoft's approach to agentic AI is about **pairing innovation with pragmatism**. They urge customers to leverage AI agents to *bridge the gap between knowledge and action*, unlocking new efficiencies in everything from customer service to finance, **but to do so incrementally and with guardrails**<sup>132 133</sup>. By providing a spectrum of tools - from no-code Copilot experiences to the Azure AI Foundry platform - Microsoft gives organizations a roadmap to experiment, validate, and scale agentic AI solutions across all industries. Crucially, this roadmap is underpinned by **responsible AI practices**: comprehensive governance frameworks, transparency through monitoring and disclosure, and diligent alignment to each organization's goals and policies. The result is an agentic AI strategy where companies can enjoy the benefits of autonomous, goal-driven AI (like *faster processes, improved service, and new insights*) **while staying in control** of outcomes. Microsoft's recommendations to customers encapsulate a simple mantra: **start small, think big, and stay responsible** - start with quick wins, plan for broader transformation, and ensure every AI agent is trustworthy, accountable, and serving your business's purpose from day one<sup>134 135</sup>.